


Limited-data learning in music

Christos Plachouras

2023



Finding

Is in-domain audio representation learning feasible? Current representation learning models aren't as magical of a solution as they sometimes seem to be; when randomly initialized without pretraining, they often produce representations that perform close to state-of-the-art. With a few minutes of training data, the gap can close further.

Question

Currently, limited-data scenarios, particularly in a continual learning settings, are tackled with scenario-specific methods. 1. Could we understand what data and task characteristics (e.g. weakly-labeled, many increments etc.) lead to different behaviors (plasticity, forgetting), and 2. is it possible to develop more universal approaches for music?