


Disentangled Representation Learning for Music Audio

Yin-Jyun Luo

2020



Finding

Exploiting general inductive biases, such as the difference in temporal resolution of factors of variation, facilitates learning semantically meaningful feature representations without supervision.

Question

Beyond controllable audio generation, does learning semantically meaningful or disentangled representations benefit classic MIR tasks such as music source separation? And how to apply the technique to more complex data such as music mixture?